

EECS 470

Power and Architecture

Many slides taken from Prof. David Brooks, Harvard University and modified by Mark Brehob . A couple of slides are also taken from Prof. Wenisch. Any errors are almost certainly Mark's.

Thanks to both!

Outline

- **Why is power a problem?**
- **What uses power in a chip?**
- **How can we reduce power?**

Outline


- Why is power a problem?
- What uses power in a chip?
- Relationship between power and performance.
- How can we reduce power?

Why is power a problem in a μ P?

- **Power used by the μ P, vs. system power**
- **Dissipating Heat**
 - Melting (very bad)
 - Packaging (to cool \rightarrow \$)
 - Heat leads to poorer performance.
- **Providing Energy**
 - Battery
 - Cost of electricity


Why worry about power dissipation?

Battery
life



Thermal issues: affect
cooling, packaging,
reliability, timing


Environment



Where does power go?

- Obviously desktop and servers are going to be different.
 - But if CPUs are a small fraction of the power, maybe we don't care much?
- Let's take a look.
 - But first a few caveats
 - I can't find recent studies on this
 - I do find that different studies can get radically different results. I'm using some fairly well-discussed numbers.


Where does the juice go in laptops?




- **Microsoft, 2009.**
- **The processor power can be a lot higher depending on what the laptop is doing. [Hsu+Kremer, 2002]**

What about servers?

SunFire T2000



Need whole-system approaches to save energy

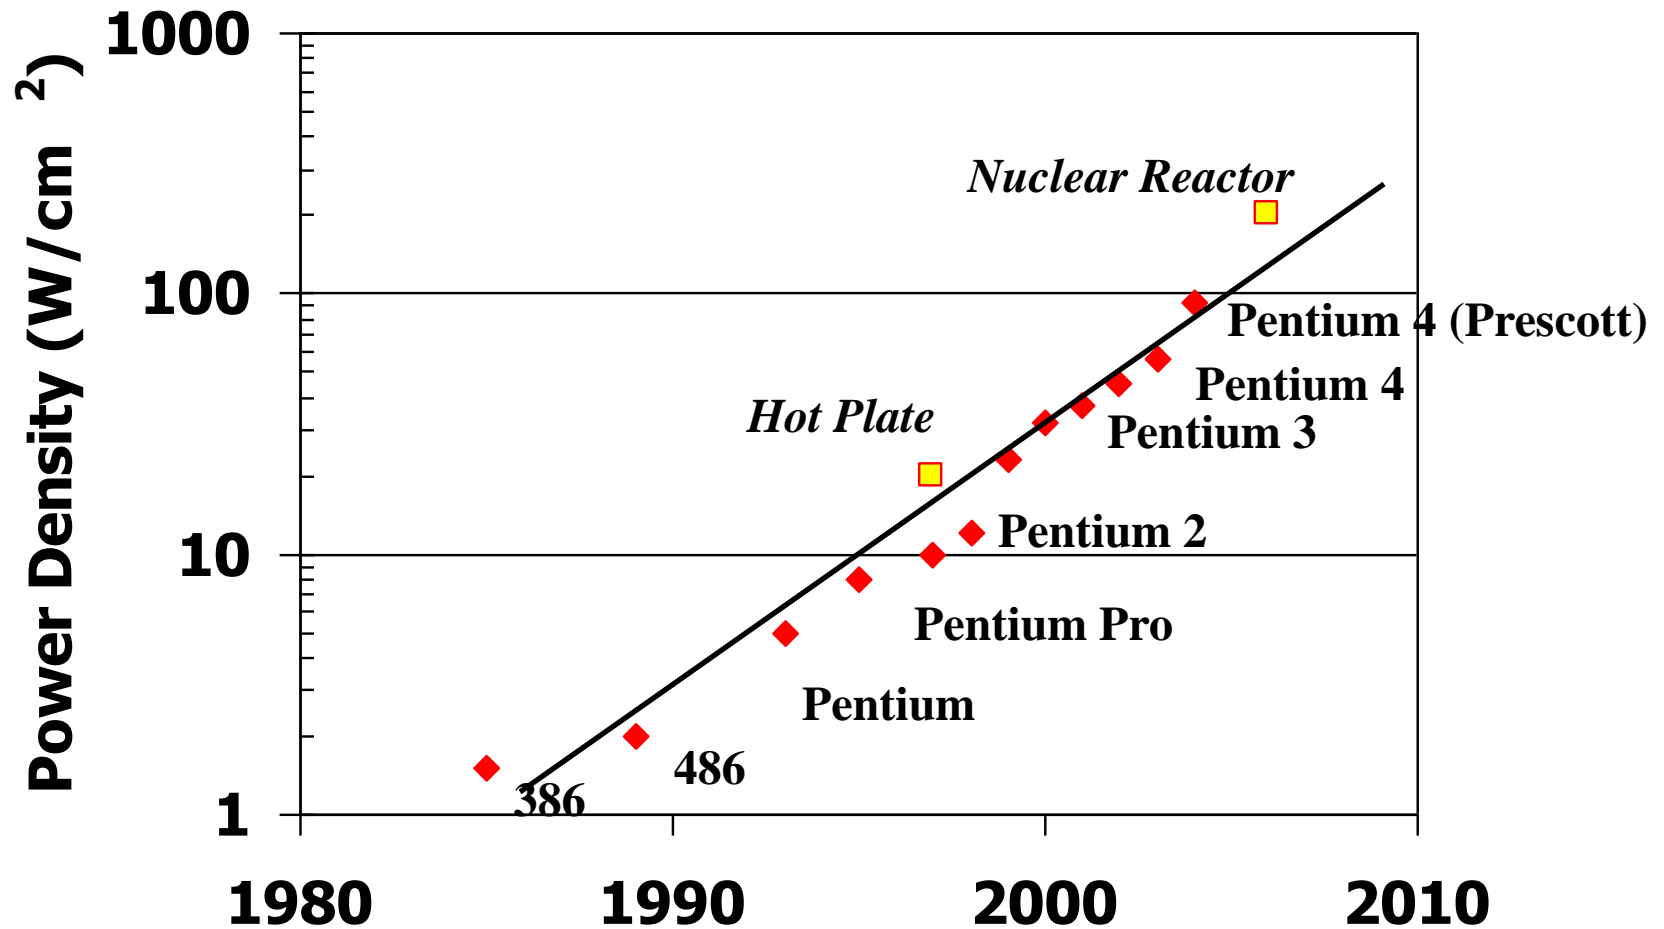


Power usage effectiveness (PUE)


A PUE of 2.0 means that for every 2W of power supplied to the data center, only 1W is consumed by computing equipment. Values of around 2.9 are pretty common

Total Power Dissipation Trends

Why is power a problem?



A Paradigm Shift In Computing



IEEE Computer—April 2001
T. Mudge

Limits on heat extraction

Stagnates performance growth

Limits on energy-efficiency of operations


Era of High Performance Computing

c. 2000


Era of Energy-Efficient Computing

Spot Heat Issues in Microprocessors

Why is power a problem?



Data center energy use




Source: US EPA 2007—Newest I can find

In 2012,
~2.0% of US energy
~\$7 billion/yr.
(Rich Miller, 2012)


Source: Mankoff et al, IEEE Computer 2008

0.5% of world CO₂ emissions;
rivals entire Czech Republic



Improving energy efficiency is a critical challenge

Where does all the power go?




Servers account for barely half of power

- 1W of cooling per 1.5W of IT load
- 10MW data center: cooling costs \$4M to \$8M / yr.

System designers must think about cooling

This may be getting worse.



- **UPS:** Power supplies/converters
- **CRAC:** Computer Room Air Conditioner.
- **PDU:** Power distribution unit

Power Usage Effectiveness (PUE)

A PUE of 2.0 means that for every 2W of power supplied to the data center, only 1W is consumed by computing equipment. Values of around 2.8 are pretty common these days, while 1.9 was a common number in 2007.


- Not clear why.

<http://ecmweb.com/energy-efficiency/data-center-efficiency-trends>

Why is cooling so costly? (1)

Server density increasing

- Integration, disaggregation reduce hardware costs
- Need for high-BW interconnect
- Data center floor space costs up to \$15,000 /m²




Heat flux up to
500W per ft² floor space;
Racks draw 4 to 20kW each

Why is cooling so costly? (2)

Heat density drives cooling cost

Cooling power grows super-linearly with thermal load




Source: C. Patel, HP Labs

*Servers' 3-year power & cooling costs
nearing their purchase price*

Intel Itanium packaging

Complex and expensive (note heatpipe)




Why is power a problem?

Source: H. Xie et al. "Packaging the Itanium Microprocessor"
Electronic Components and Technology Conference 2002


P4 packaging

Why is power a problem?

- **Simpler, but still...**



From Tiwari, et al., DAC98



Source: Intel web site

Power-Aware Computing Applications

↑
Temperature/dt-Constrained
—



— Energy-Constrained Computing —→

Environment

Why is power a problem?

- **Environment Protection Agency (EPA): computers consume 10% of commercial electricity consumption**
 - This incl. peripherals, possibly also manufacturing
 - A DOE report suggested this percentage is much lower (3.0-3.5%)
 - No consensus, but it's still a lot
 - Interesting to look at the numbers:
 - <http://enduse.lbl.gov/projects/infotech.html>
- **Data center growth was cited as a contribution to the 2000/2001 California Energy Crisis**
- **Equivalent power (with only 30% efficiency) for AC**
- **CFCs used for refrigeration**
- **Lap burn**
- **Fan noise**

Power-Aware Needed across all computing platforms



Why is power a problem?

- **Mobile/portable (cell phones, laptops, PDA)**
 - Battery life is critical
- **Desktops/Set-Top (PCs and game machines)**
 - Packaging cost is critical
- **Servers (Mainframes and compute-farms)**
 - Packaging limits
 - Volumetric (performance density)

What uses power in a chip?



How CMOS Transistors Work

What uses power in a chip?




MOS Transistors are Switches

What uses power in a chip?



Static CMOS




What uses power in a chip?



PUN and PDN are Dual Networks

Basic Logic Gates

What uses power in a chip?



CMOS Water Analogy

What uses power in a chip?

Electron: water molecule

Charge: weight of water



Voltage: height

Current: flow rate

Capacitance: container cross-section

(Think of power-plants that store energy in water towers)

Liquid Inverter



- **Capacitance at input**
 - **Gates of NMOS, PMOS**
 - **Metal interconnect**
- **Capacitance at output**
 - **Fanout (# connections) to other gates**
 - **Metal Interconnect**

NMOS conducts when water level is above switching threshold


PMOS conducts below

No conduction after container full


Slide courtesy D. Brooks, Harvard

Inverter Signal Propagation (1)

$t < 0$
 $V_{in} = 0$




$t = 0$
 $V_{in} = 1$




Slide courtesy D. Brooks, Harvard

Inverter Signal Propagation (2)

$t = 1$
 $V_{in} = 0$



$t = 2$
 $V_{in} = 1$



Slide courtesy D. Brooks, Harvard

Delay and Power Observations

What uses power in a chip?


- **Load capacitance increases delay**
 - High fanout (gates attached to output)
 - Interconnection
- **Higher current can increase speed**
 - Increasing transistor width raises currents but also raises capacitance
- **Energy per switching event independent of current**
 - Depends on amount of charge moved, not rate

Power: The Basics

What uses power in a chip?

- **Dynamic power vs. Static power**
 - **Dynamic: “switching” power**
 - **Static: “leakage” power**
 - **Dynamic power dominates, but static power increasing in importance**
- **Static power: steady, per-cycle energy cost**
- **Dynamic power: capacitive and short-circuit**
- **Capacitive power: charging/discharging at transitions from $0 \rightarrow 1$ and $1 \rightarrow 0$**
- **Short-circuit power: power due to brief short-circuit current during transitions.**

Dynamic (Capacitive) Power Dissipation



- Data dependent – a function of **switching** activity

Capacitive Power dissipation

What uses power in a chip?

Capacitance:
Function of wire
length, transistor size

Supply Voltage:
Has been dropping
with successive fab
generations

$$\text{Power} \sim \frac{1}{2} CV^2Af$$

Activity factor:
How often, on average,
do wires switch?


Clock frequency:
Increasing...

Lowering Dynamic Power

What uses power in a chip?

- **Reducing V_{dd} has a quadratic effect**
 - Has a negative (\sim linear) effect on performance however
- **Lowering C_L**
 - May improve performance as well
 - Keep transistors small (keeps intrinsic capacitance (gate and diffusion) small)
- **Reduce switching activity**
 - A function of signal transition stats and clock rate
 - Clock Gating idle units
 - Impacted by logic and architecture decisions

Static Power: Leakage Currents



$$I_{DSub} = k \cdot e^{\frac{-q \cdot V_T}{a \cdot k_a \cdot T}}$$

- Subthreshold currents grow exponentially with increases in temperature, decreases in threshold voltage
 - But threshold voltage scaling is key to circuit performance!
- Gate leakage primarily dependent on gate oxide thickness, biases
- Both type of leakage heavily dependent on stacking and input pattern

Lowering Static Power

What uses power in a chip?

- **Design-time Decisions**
 - Use fewer, smaller transistors -- stack when possible to minimize contacts with Vdd/Gnd
 - Multithreshold process technology (multiple oxides too!)
 - Use “high-Vt” slow transistors whenever possible
- **Dynamic Techniques**
 - Reverse-Body Bias (dynamically adjust threshold)
 - Low-leakage sleep mode (maintain state), e.g. XScale
 - Vdd-gating (Cut voltage/gnd connection to circuits)
 - Zero-leakage sleep mode
 - Lose state, overheads to enable/disable



Power vs. Energy

What uses power in a chip?

- **Power** consumption in **Watts**
 - Determines battery life in hours
 - Sets packaging limits
- **Energy** efficiency in **joules**
 - Rate at which energy is consumed over time
 - $\text{Energy} = \text{power} * \text{delay}$ (joules = watts * seconds)
 - Lower energy number means less power to perform a computation at same frequency

Power vs. Energy

What uses power in a chip?



Power vs. Energy

What uses power in a chip?

- **Power-delay Product (PDP) = $P_{avg} * t$**
 - PDP is the average energy consumed per switching event
- **Energy-delay Product (EDP) = $PDP * t$**
 - Takes into account that one can trade increased delay for lower energy/operation
- **Energy-delay² Product (EDDP) = $EDP * t$**
 - Why do we need so many formulas?!?
 - We want a voltage-invariant efficiency metric! Why?
 - Power $\sim \frac{1}{2} CV^2Af$, Performance $\sim f$ (and V)

E vs. EDP vs. ED²P

What uses power in a chip?

- **Power $\sim CV^2f \sim V^3$ (fixed microarch/design)**
- **Performance $\sim f \sim V$ (fixed microarch/design)**
- **(For the nominal voltage range, f varies linearly with V)**
- **Comparing processors that can only use freq/voltage scaling as the primary method of power control:**
 - **$(\text{perf})^3 / \text{power}$, or MIPS^3 / W is a fair metric to compare energy efficiencies.**
 - **This is an ED² P metric. We could also use: $(\text{CPI})^3 * W$ for a given application**

E vs. EDP vs. ED²P

What uses power in a chip?


- **Currently have a processor design:**
 - **80W, 1 BIPS, 1.5V, 1GHz**
 - **Want to reduce power, willing to lose some performance**
 - **Cache Optimization:**
 - **IPC decreases by 10%, reduces power by 20% =>**
Final Processor: 900 MIPS, 64W
 - **Relative E = MIPS/W (higher is better) =**
14/12.5 = 1.125x
 - **Energy is better, but is this a “better” processor?**

Not necessarily

- **80W, 1 BIPS, 1.5V, 1GHz**
 - **Cache Optimization:**
 - IPC decreases by 10%, reduces power by 20% => Final Processor: 900 MIPS, 64W
 - Relative E = MIPS/W (higher is better) = $14/12.5 = 1.125x$
 - Relative EDP = $\text{MIPS}^2/\text{W} = 1.01x$
 - Relative ED²P = $\text{MIPS}^3/\text{W} = .911x$
- **What if we just adjust frequency/voltage on processor?**
 - How to reduce power by 20%?
 - $P = CV^2F = CV^3 \Rightarrow$ Drop voltage by 7% (and also Freq) => $.93 \cdot .93 \cdot .93 = .8x$
 - So for equal power (64W)
 - Cache Optimization = 900MIPS
 - Simple Voltage/Frequency Scaling = 930MIPS

What do we mean by Power?

What uses power in a chip?



- **Max Power:** Artificial code generating max CPU activity
- **Worst-case App Trace:** *Practical* applications worst-case
- **Thermal Power:** Running average of worst-case app power over a time period corresponding to thermal time constant
- **Average Power:** Long-term average of typical apps (minutes)
- **Transient Power:** Variability in power consumption for supply net